

# Rock vs Hip-Hop & Pop

*Naga Santhosh Kartheek Karnati*

*8/23/2020*

## R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
#Loading libraries
```

```
library(ggplot2)
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library(rvest)
```

```
## Loading required package: xml2
```

```
library(xml2)
library(selectr)
library(stringr)
library(jsonlite)
library(naniar)
library(tokenizers)
library(tidytext)
library(dbplyr)
```

```
## Warning: package 'dbplyr' was built under R version 3.6.3
```

```
##
```

```
## Attaching package: 'dbplyr'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      ident, sql
```

```

library(tidyr)
library(tibble)
library(data.table)

##
## Attaching package: 'data.table'

## The following objects are masked from 'package:dplyr':
##
##      between, first, last

#Loading dataframes

pewords <- fread("D:/NEU/Individual Projects/Bands project/Data/popsong_ele_words.csv")
rewords <- fread("D:/NEU/Individual Projects/Bands project/Data/song_lyrics_ele_words.csv")
pwords <- fread("D:/NEU/Individual Projects/Bands project/Data/popsong_words.csv")
rwords <- fread("D:/NEU/Individual Projects/Bands project/Data/song_lyrics_words.csv")

# View(pewords)
# View(rewords)
# View(pwords)
# View(rwords)

#Merging dataframes

ewords <- rbind(rewords, pewords, use.names = FALSE)
words <- rbind(rwords, pwords, use.names = FALSE)

# View(ewords)
# View(words)

#Save dataframes

#ewords df
write.csv(ewords,
          "D:/NEU/Individual Projects/Bands project/Data/ewords.csv",
          row.names = FALSE)

#words df
write.csv(words,
          "D:/NEU/Individual Projects/Bands project/Data/words.csv",
          row.names = FALSE)

#Comparing vocabulary

n_words <- words %>% count(band, word, sort = TRUE)
n_words

## Warning: `...` is not empty.
##
## We detected these problematic arguments:

```

```
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 39,573 x 3
##   band          word      n
##   <chr>         <chr> <int>
## 1 The Rolling Stones baby    712
## 2 The Rolling Stones yeah    550
## 3 The Rolling Stones love    484
## 4 Queen          love    459
## 5 Chris Brown     girl    433
## 6 Aerosmith       yeah    377
## 7 Queen          yeah    365
## 8 Aerosmith       love    360
## 9 Rihanna         love    345
## 10 AC/DC          rock    330
## # ... with 39,563 more rows
```

```
total_words <- n_words %>% group_by(band) %>% summarise(words = sum(n))
total_words
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 20 x 2
##   band          words
##   <chr>         <int>
## 1 AC/DC          11255
## 2 Adele           2393
## 3 Aerosmith       17226
## 4 Chris Brown     12000
## 5 Eagles           5224
## 6 Guns n Roses     7611
## 7 John Legend      3142
## 8 Led Zeppelin     5548
## 9 Metallica        11200
## 10 Miley Cyrus      3115
## 11 One Direction    7421
## 12 Pink            10364
## 13 Pink Floyd       8354
## 14 Queen           15877
## 15 Rihanna          11617
## 16 Selena Gomez      7588
## 17 Shawn Mendes      391
## 18 The Rolling Stones 24025
## 19 The Who          12425
## 20 Travis Scott     313
```

```

unique_words <- n_words %>% select(-n) %>% count(band, sort = TRUE)

vocabulary <- total_words %>% inner_join(unique_words) %>% mutate(ratio=n/words) %>% arrange(desc(ratio))

## Joining, by = "band"

vocabulary

## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?

## # A tibble: 20 x 4
##   band      words    n ratio
##   <chr>    <int> <int> <dbl>
## 1 Travis Scott    313   202 0.645
## 2 Shawn Mendes   391   151 0.386
## 3 John Legend   3142  1032 0.328
## 4 Pink Floyd    8354  2681 0.321
## 5 Eagles       5224  1612 0.309
## 6 Miley Cyrus   3115   945 0.303
## 7 Adele        2393   712 0.298
## 8 Led Zeppelin   5548  1519 0.274
## 9 Guns n Roses   7611  2055 0.270
## 10 The Who      12425  3160 0.254
## 11 Metallica    11200  2725 0.243
## 12 Pink        10364  2237 0.216
## 13 Aerosmith    17226  3632 0.211
## 14 AC/DC        11255  2330 0.207
## 15 One Direction 7421  1477 0.199
## 16 Queen        15877  3137 0.198
## 17 Chris Brown  12000  2340 0.195
## 18 Selena Gomez   7588  1426 0.188
## 19 The Rolling Stones 24025  4297 0.179
## 20 Rihanna      11617  1903 0.164

#Ignore Travis Scott and Shawn mendes since we have less data .
#John legend most
#Rihanna least

#Comparing rock bands vs hip-hop and pop singers
bands <- c("Pink Floyd", "Eagles", "Led Zeppelin", "Guns n Roses", "The Who", "Metallica", "Aerosmith",
           "AC/DC", "Queen", "The Rolling Stones")

singers <- c("Travis Scott", "Shawn Mendes", "John Legend", "Miley Cyrus", "Adele", "Pink", "One Direction",
             "Chris Brown", "Selena Gomez", "Rihanna")

vocabulary %>% filter(band %in% bands) %>% summarise(no_rockwords = sum(words),
                                                    unique_rockwords = sum(n)) %>% mutate(ratio=unique_rockwords/no_rockwords)

```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 1 x 3
##   no_rockwords unique_rockwords ratio
##       <int>           <int> <dbl>
## 1       118745           27148 0.229
```

```
vocabulary %>% filter(band %in% singers) %>% summarise(no_hhpwords = sum(words),
  unique_hhpwords = sum(n)) %>% mutate(ratio=unique_hhpwords/no_hhpwords)
```

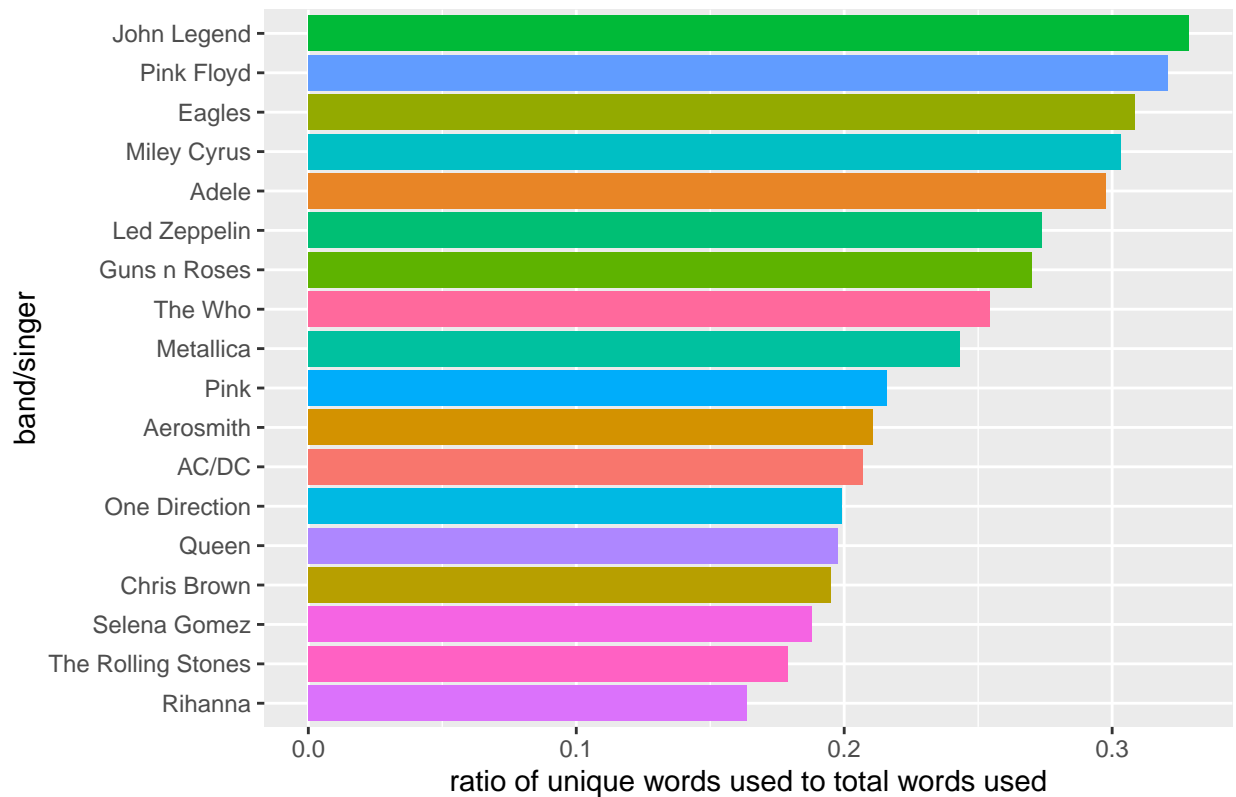
```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 1 x 3
##   no_hhpwords unique_hhpwords ratio
##       <int>           <int> <dbl>
## 1       58344           12425 0.213
```

```
#rock bands have better vocabulary than hip-hop and pop artists.
```

```
vocabulary %>% filter(!band %in% c("Travis Scott", "Shawn Mendes")) %>%
  ggplot(aes(reorder(band, ratio), ratio, fill = band)) + geom_col(show.legend = FALSE) +
  labs(x = "band/singer", y="ratio of unique words used to total words used",
    title = "Comparison of vocabulary") + coord_flip()
```

## Comparison of vocabulary



#Topic modeling after eliminating words with low tf\_idf scores

```
library(tokenizers)
library(tidytext)
library(tidyr)
library(tm)
```

```
## Warning: package 'tm' was built under R version 3.6.3
```

```
## Loading required package: NLP
```

```
##
```

```
## Attaching package: 'NLP'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## annotate
```

```
library(topicmodels)
```

#Term frequencies

```
##Calculating the term frequency per each band/singer:
```

```
tfs <- n_words %>%  
  left_join(total_words)%>%  
  mutate(tf = n/words) %>%  
  arrange(desc(tf))
```

```
## Joining, by = "band"
```

```
tfs
```

```
## Warning: `...` is not empty.
```

```
##
```

```
## We detected these problematic arguments:
```

```
## * `needs_dots`
```

```
##
```

```
## These dots only exist to allow future extensions and should be empty.
```

```
## Did you misspecify an argument?
```

```
## # A tibble: 39,573 x 5
```

```
##   band      word      n words      tf  
##   <chr>      <chr> <int> <int> <dbl>  
## 1 Travis Scott yeah      43    313 0.137  
## 2 Adele      love     111   2393 0.0464  
## 3 Shawn Mendes mercy     17    391 0.0435  
## 4 Selena Gomez love     306   7588 0.0403  
## 5 John Legend love     126   3142 0.0401  
## 6 Chris Brown girl     433  12000 0.0361  
## 7 Shawn Mendes baby      13    391 0.0332  
## 8 Travis Scott light     10    313 0.0319  
## 9 One Direction na       237   7421 0.0319  
## 10 Led Zeppelin baby     176   5548 0.0317  
## # ... with 39,563 more rows
```

```
#tf_idfs
```

```
unreq_words <- tibble(word = c("chorus","chris", "53rd", "miley", "pink", "rihanna"))
```

```
#tf_idfs
```

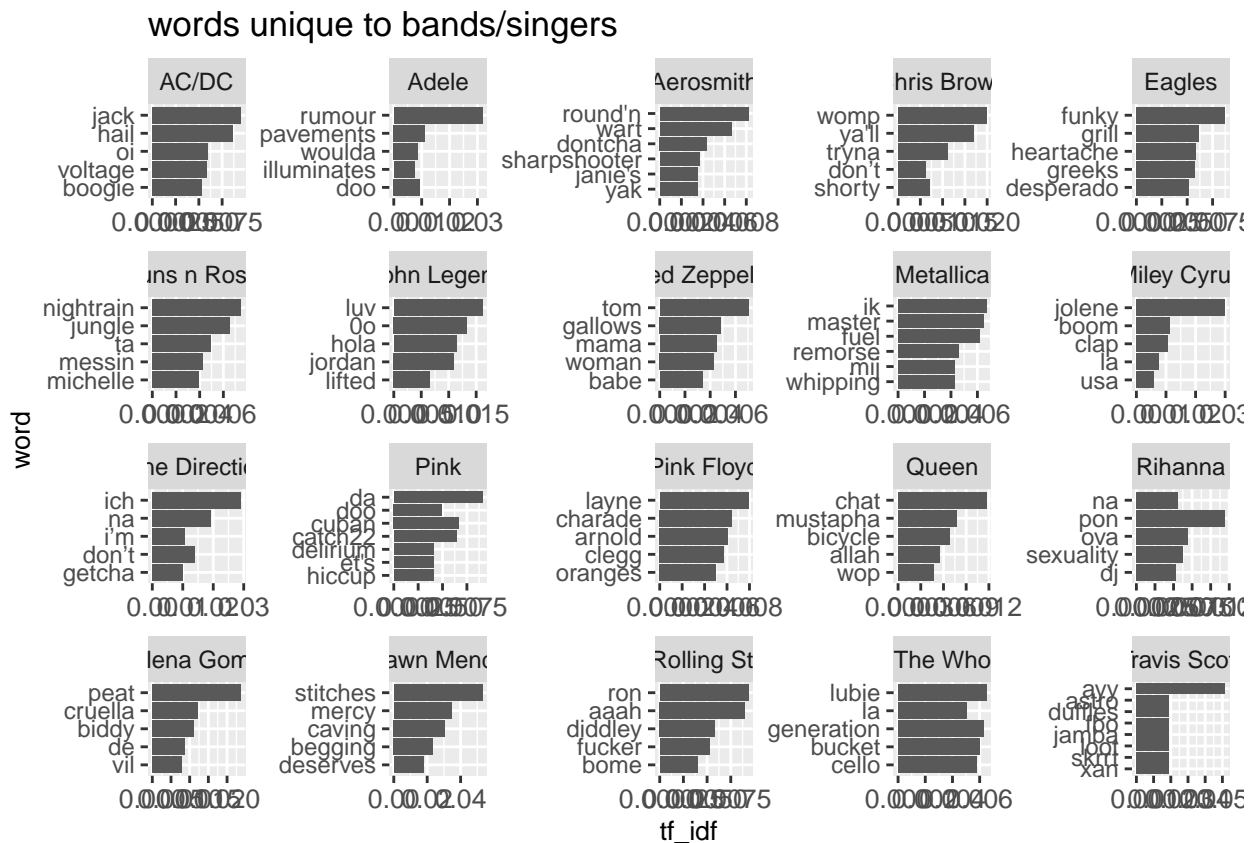
```
tf_idfs <- words %>%  
  anti_join(unreq_words) %>%  
  count(band, word, sort=TRUE) %>%  
  bind_tf_idf(word, band, n) %>%  
  arrange(desc(tf_idf))
```

```
## Joining, by = "word"
```

```
My_Theme = theme(  
  axis.title.x = element_text(size = 10),  
  axis.text.x = element_text(size = 10),  
  axis.title.y = element_text(size = 10))
```

```
tf_idfs %>%
  arrange(desc(tf_idf)) %>%
  mutate(word = factor(word, levels = rev(unique(word)))) %>%
  group_by(band) %>%
  top_n(5) %>%
  ungroup() %>%
  ggplot(aes(reorder(word,tf_idf), tf_idf), fill = band) +
  geom_col(show.legend = FALSE) +
  labs(x = NULL, y = "tf-idf") +
  facet_wrap(~band, ncol = 5, scales = "free") +
  labs(x="word", y="tf_idf", title = "words unique to bands/singers")+
  coord_flip() + My_Theme
```

```
## Selecting by tf_idf
```



```
tf_idfs_top_5 <- tf_idfs %>%
  arrange(desc(tf_idf)) %>%
  mutate(word = factor(word, levels = rev(unique(word)))) %>%
  group_by(band) %>%
  top_n(5) %>%
  ungroup()
```

```
## Selecting by tf_idf
```



```
summary(tf_idfs$tf_idf)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## 0.0000000 0.0001154 0.0001992 0.0003820 0.0003643 0.0530007
```

```
#remove the bottom 25% of words from the tf_idf df
```

```
#Remove words in the bottom 25% of tf_idf scores
```

```
tf_idfs_filtered <- tf_idfs %>% filter(tf_idf > 0.0001154) %>% select(-n,-tf,-idf,-tf_idf)
#dim(tf_idfs)
```

```
#join words df and tf_idfs_filtered df on composite key (band and word)
```

```
words <- words %>% inner_join(tf_idfs_filtered, by = c("band"="band", "word"="word"))
```

```
words$genre[words$band %in% bands] <- "Rock"
```

```
words$genre[words$band %in% singers] <- "Hip-Hop & Pop"
```

```
#Words unique to rock bands and hip-hop and pop artists
```

```
words %>% unique() %>% group_by(genre) %>% top_n(5, tf_idf) %>% ungroup()
```

```
## Warning: `...` is not empty.
```

```
##
```

```
## We detected these problematic arguments:
```

```
## * `needs_dots`
```

```
##
```

```
## These dots only exist to allow future extensions and should be empty.
```

```
## Did you misspecify an argument?
```

```
## # A tibble: 10 x 7
```

|       | band               | word     | n     | tf      | idf   | tf_idf  | genre         |
|-------|--------------------|----------|-------|---------|-------|---------|---------------|
|       | <chr>              | <chr>    | <int> | <dbl>   | <dbl> | <dbl>   | <chr>         |
| ## 1  | The Rolling Stones | ron      | 75    | 0.00312 | 3.00  | 0.00935 | Rock          |
| ## 2  | The Rolling Stones | aaah     | 72    | 0.00300 | 3.00  | 0.00898 | Rock          |
| ## 3  | Eagles             | funky    | 24    | 0.00460 | 1.90  | 0.00872 | Rock          |
| ## 4  | AC/DC              | jack     | 179   | 0.0159  | 0.598 | 0.00951 | Rock          |
| ## 5  | Queen              | chat     | 62    | 0.00391 | 3.00  | 0.0117  | Rock          |
| ## 6  | Adele              | rumour   | 33    | 0.0138  | 2.30  | 0.0318  | Hip-Hop & Pop |
| ## 7  | Travis Scott       | ayy      | 7     | 0.0224  | 2.30  | 0.0515  | Hip-Hop & Pop |
| ## 8  | Shawn Mendes       | caving   | 4     | 0.0102  | 3.00  | 0.0306  | Hip-Hop & Pop |
| ## 9  | Shawn Mendes       | mercy    | 17    | 0.0435  | 0.799 | 0.0347  | Hip-Hop & Pop |
| ## 10 | Shawn Mendes       | stitches | 9     | 0.0230  | 2.30  | 0.0530  | Hip-Hop & Pop |

```
#Topic modeling with each band/singer as a document
```

```
words2 <- words %>% select(band, word)
```

```
#DTM with each band/singer as a document
```

```
bandsinger_dtm <- words2 %>%
  count(band, word) %>%
  cast_dtm(document = band, term = word, value = n)

bandsinger_dtm
```

```
## <<DocumentTermMatrix (documents: 20, terms: 14363)>>
## Non-/sparse entries: 29738/257522
## Sparsity          : 90%
## Maximal term length: 99
## Weighting         : term frequency (tf)
```

```
lda1 <- LDA(bandsinger_dtm, k = 20, control = list(seed = 5))
lda1
```

```
## A LDA_VEM topic model with 20 topics.
```

```
topics1 <- tidy(lda1, matrix = "beta")
topics1
```

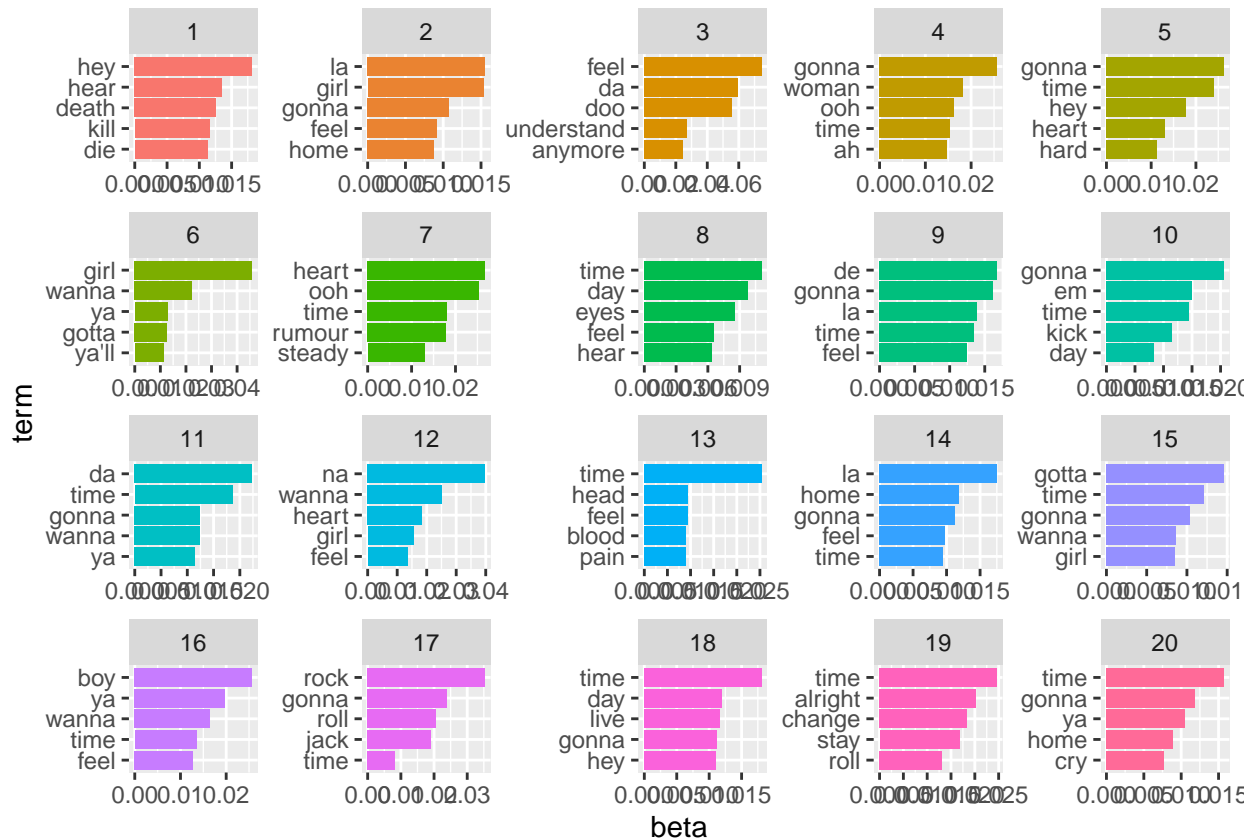
```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 287,260 x 3
##   topic term      beta
##   <int> <chr>   <dbl>
## 1     1 11x 1.03e-240
## 2     2 11x 5.19e-242
## 3     3 11x 2.34e-238
## 4     4 11x 1.21e-240
## 5     5 11x 2.20e-241
## 6     6 11x 6.60e-242
## 7     7 11x 4.75e-239
## 8     8 11x 2.13e-241
## 9     9 11x 8.32e-242
## 10    10 11x 1.10e-240
## # ... with 287,250 more rows
```

```
topterms1 <- topics1 %>%
  group_by(topic) %>%
  top_n(5, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)

topterms1 %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
```

```
ggplot(aes(term, beta, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  coord_flip() +
  scale_x_reordered()
```



*#which topics are associated with each singer (a document)?*

```
gamma1 <- tidy(lda1, matrix = "gamma")
gamma1
```

```
## Warning: `...` is not empty.
```

```
##
```

```
## We detected these problematic arguments:
```

```
## * `needs_dots`
```

```
##
```

```
## These dots only exist to allow future extensions and should be empty.
```


```
## Did you misspecify an argument?
```

```
## # A tibble: 400 x 3
```

```
##   document    topic    gamma
##   <chr>      <int>    <dbl>
## 1 AC/DC        1 0.000000906
## 2 Adele        1 0.00000412
## 3 Aerosmith    1 0.00000642
```

```
## 4 Chris Brown      1 0.000000898
## 5 Eagles           1 0.00000191
## 6 Guns n Roses     1 0.00000133
## 7 John Legend      1 0.00000333
## 8 Led Zeppelin     1 0.00000185
## 9 Metallica        1 0.554
## 10 Miley Cyrus     1 0.00000312
## # ... with 390 more rows
```

```
gamma1 %>%
  mutate(document = reorder(document, gamma * topic)) %>%
  ggplot(aes(factor(topic), gamma)) +
  geom_boxplot() +
  facet_wrap(~ document, ncol = 4)
```



```
#The rolling stones associated with topics 2 and 5
#Metallica associated with topics 1 and 13
#Adele associated with topics 7 and 3
#Travis scott associated with topics 10 and 6
#Pink associated with topics 11 and 3
#Shawn mendes associated with topics 13, 12, 19, 12, 7 and 9
```

```
#band/singer classifications
```

```
bandsinger_classifications <- gamma1 %>%
  group_by(document) %>%
  top_n(1, gamma) %>%
  ungroup()
```

```
bandsinger_classifications
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 20 x 3
##   document      topic gamma
##   <chr>      <int> <dbl>
## 1 Metallica      1 0.554
## 2 The Rolling Stones  2 0.589
## 3 Led Zeppelin     4 1.00
## 4 Chris Brown     6 1.00
## 5 Adele           7 0.880
## 6 Pink Floyd       8 1.00
## 7 Miley Cyrus     9 1.00
## 8 Selena Gomez     9 1.00
## 9 Eagles          10 1.00
## 10 Travis Scott    10 0.901
## 11 Pink            11 0.892
## 12 One Direction   12 1.00
## 13 Shawn Mendes    13 0.560
## 14 The Who         14 0.999
## 15 Aerosmith       15 1.00
## 16 Rihanna         16 1.00
## 17 AC/DC           17 1.00
## 18 Queen           18 1.00
## 19 John Legend     19 1.00
## 20 Guns n Roses    20 1.00
```

```
#assigning words in each document to a topic
```

```
bandsinger_assignments <- augment(lda1, data = bandsinger_dtm)
bandsinger_assignments
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 29,738 x 4
##   document      term count .topic
##   <chr>         <chr> <dbl> <dbl>
## 1 AC/DC         11x      1     17
## 2 AC/DC         15      1     17
## 3 Adele         15      1      7
## 4 Chris Brown   15      1      6
## 5 Guns n Roses  15      1     20
## 6 AC/DC         24      1     17
## 7 Miley Cyrus   24      1      9
## 8 Selena Gomez  24      2      9
## 9 The Rolling Stones 24      2      2
## 10 AC/DC        3.00     1     17
## # ... with 29,728 more rows
```

```
bandsinger_assignments <- bandsinger_assignments %>%
  inner_join(bandsinger_classifications, by = c(".topic" = "topic"))
bandsinger_assignments
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

```
## # A tibble: 32,317 x 6
##   document.x term count .topic document.y gamma
##   <chr>      <chr> <dbl> <dbl> <chr>      <dbl>
## 1 AC/DC      11x      1     17 AC/DC      1.00
## 2 AC/DC      15      1     17 AC/DC      1.00
## 3 Adele      15      1      7 Adele      0.880
## 4 Chris Brown 15      1      6 Chris Brown 1.00
## 5 Guns n Roses 15      1     20 Guns n Roses 1.00
## 6 AC/DC      24      1     17 AC/DC      1.00
## 7 Miley Cyrus 24      1      9 Miley Cyrus 1.00
## 8 Miley Cyrus 24      1      9 Selena Gomez 1.00
## 9 Selena Gomez 24      2      9 Miley Cyrus 1.00
## 10 Selena Gomez 24      2      9 Selena Gomez 1.00
## # ... with 32,307 more rows
```

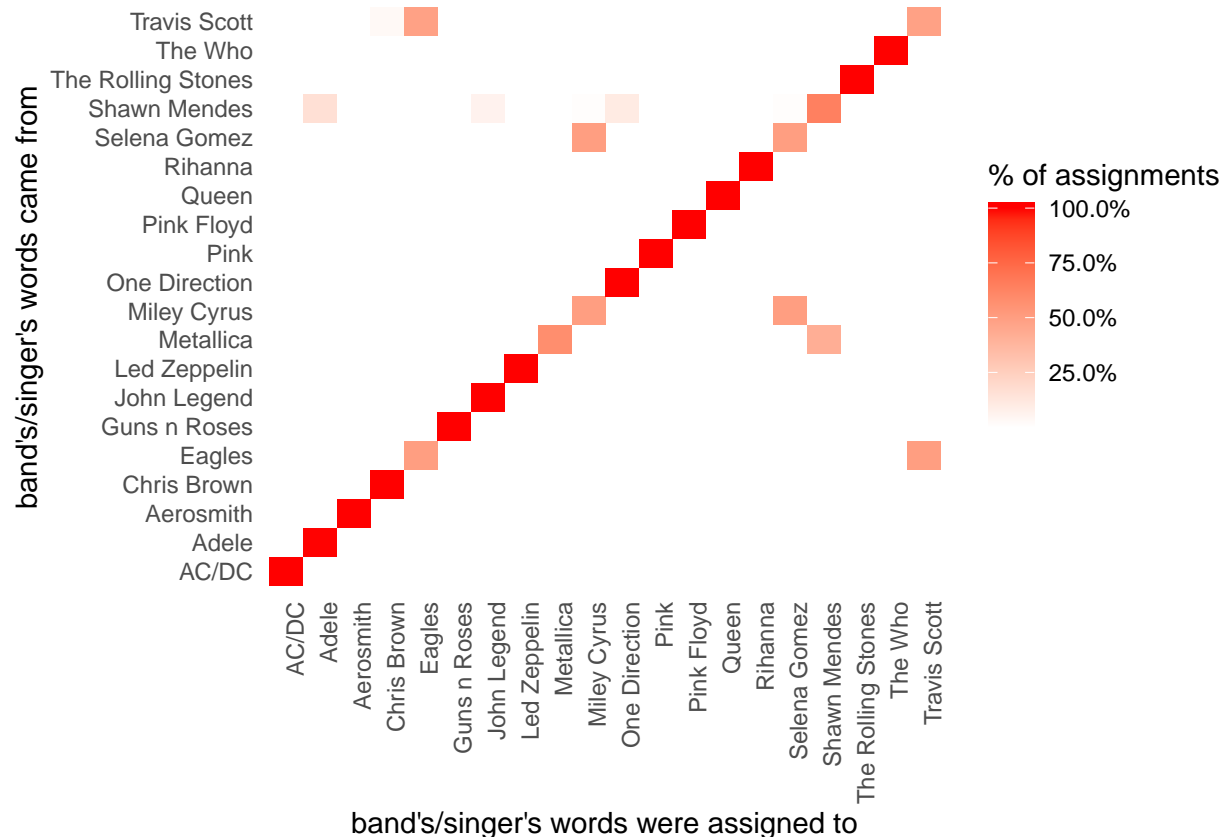
*#here document.x is the "true" band whereas document.y is the "consensus" band.*

#confusion matrix for LDA where each band/singer is a document

```
library(scales)

bandsinger_assignments %>%
  count(document.x, document.y, wt = count) %>%
  group_by(document.x) %>%
  mutate(percent = n / sum(n)) %>%
```

```
ggplot(aes(document.y, document.x, fill = percent)) +
  geom_tile() +
  scale_fill_gradient2(high = "red", label = percent_format()) +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1),
        panel.grid = element_blank()) +
  labs(x = "band's/singer's words were assigned to",
       y = "band's/singer's words came from",
       fill = "% of assignments")
```



*#The few songs that shawn mendes has can be mistaken as Adele, John Legend, One direction songs*  
*#Travis scott's song can be mistaken as an Eagles song.*  
*#Selena Gomez's songs can be mistaken as Miley cyrus songs and vice versa*

#topic modeling with a different approach

```
female_artists <- c("Adele", "Pink", "Selena Gomez", "Miley Cyrus", "Rihanna")
male_artists <- c("John Legend", "Travis Scott", "Shawn Mendes", "One Direction", "Chris Brown")

words$genre[words$band %in% male_artists] <- "Hip-Hop/Pop Male"
words$genre[words$band %in% female_artists] <- "Hip-Hop/Pop Female"

words3 <- words %>% select(genre, word)
```

#DTM with each genre as a document

```
genre_dtm <- words3 %>%
  count(genre, word) %>%
  cast_dtm(document = genre, term = word, value = n)

genre_dtm
```

```
## <<DocumentTermMatrix (documents: 3, terms: 14363)>>
## Non-/sparse entries: 19108/23981
## Sparsity          : 56%
## Maximal term length: 99
## Weighting          : term frequency (tf)
```

```
lda2 <- LDA(genre_dtm, k = 3, control = list(seed = 10))
lda2
```

```
## A LDA_VEM topic model with 3 topics.
```

```
topics2 <- tidy(lda2, matrix = "beta")
topics2
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```


```
## # A tibble: 43,089 x 3
##   topic term      beta
##   <int> <chr>   <dbl>
## 1     1  1 0    1.70e-22
## 2     2  2 0    1.38e-43
## 3     3  3 0    8.53e- 5
## 4     4  1 1    2.25e- 4
## 5     5  2 1    9.49e- 5
## 6     6  3 1    2.03e- 3
## 7     7  1 1'2   6.77e-23
## 8     8  2 1'2   2.62e-44
## 9     9  3 1'2   4.26e- 5
## 10    10  1 10    8.08e-25
## # ... with 43,079 more rows
```

```
topterms2 <- topics2 %>%
  group_by(topic) %>%
  top_n(5, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)

topterms2 %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
```



```
ggplot(aes(term, beta, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  coord_flip() +
  scale_x_reordered()
```



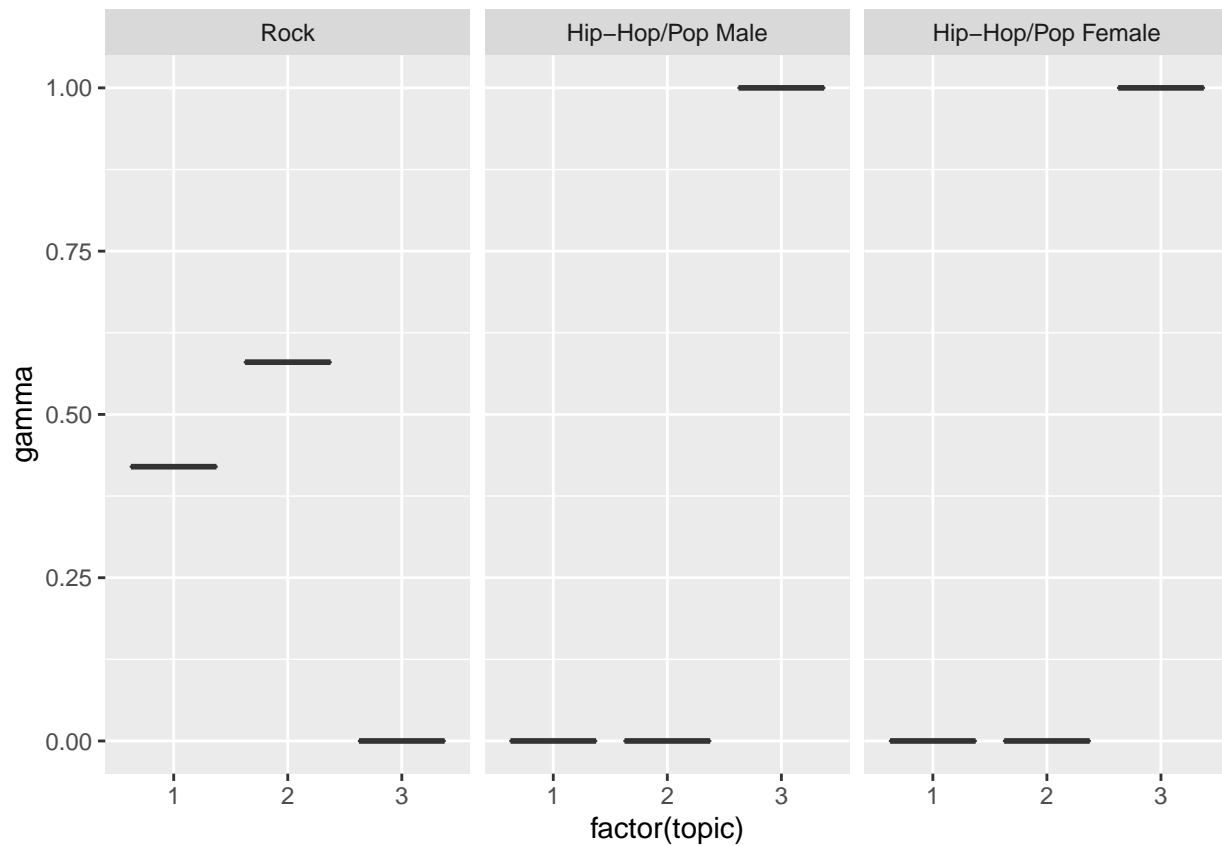
```
#which topics are associated with each genre (a document)?
gamma2 <- tidy(lda2, matrix = "gamma")
gamma2
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?

## # A tibble: 9 x 3
##   document          topic    gamma
##   <chr>             <int>   <dbl>
## 1 Hip-Hop/Pop Female     1 0.00000194
## 2 Hip-Hop/Pop Male       1 0.00000297
## 3 Rock                   1 0.420
```

```
## 4 Hip-Hop/Pop Female      2 0.00000194
## 5 Hip-Hop/Pop Male       2 0.00000297
## 6 Rock                   2 0.580
## 7 Hip-Hop/Pop Female     3 1.00
## 8 Hip-Hop/Pop Male       3 1.00
## 9 Rock                   3 0.0000325
```

```
gamma2 %>%
  mutate(document = reorder(document, gamma * topic)) %>%
  ggplot(aes(factor(topic), gamma)) +
  geom_boxplot() +
  facet_wrap(~ document, ncol = 4)
```



```
#rock bands are associated with topics 2 and 1
#hip-hop and pop artists (male and female) are associated with topic 3.
```

```
#genre classifications
```

```
genre_classifications <- gamma2 %>%
  group_by(document) %>%
  top_n(1, gamma) %>%
  ungroup()
```

```
genre_classifications
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?

## # A tibble: 3 x 3
##   document      topic gamma
##   <chr>         <int> <dbl>
## 1 Rock          2 0.580
## 2 Hip-Hop/Pop Female 3 1.00
## 3 Hip-Hop/Pop Male   3 1.00
```

#assigning words in each document to a topic

```
genre_assignments <- augment(lda2, data = genre_dtm)
genre_assignments
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?

## # A tibble: 19,108 x 4
##   document      term    count .topic
##   <chr>         <chr>   <dbl> <dbl>
## 1 Hip-Hop/Pop Female 0         4      3
## 2 Hip-Hop/Pop Female 1        66      3
## 3 Hip-Hop/Pop Male   1        29      3
## 4 Rock              1        14      1
## 5 Hip-Hop/Pop Female 1'2         2      3
## 6 Hip-Hop/Pop Female 10         2      3
## 7 Hip-Hop/Pop Male   10         3      3
## 8 Hip-Hop/Pop Female 10,000ft     2      3
## 9 Hip-Hop/Pop Female 10s         1      3
## 10 Hip-Hop/Pop Female 11         2      3
## # ... with 19,098 more rows
```

```
genre_assignments <- genre_assignments %>%
  inner_join(genre_classifications, by = c(".topic" = "topic"))
genre_assignments
```

```
## Warning: `...` is not empty.
##
## We detected these problematic arguments:
## * `needs_dots`
##
## These dots only exist to allow future extensions and should be empty.
## Did you misspecify an argument?
```

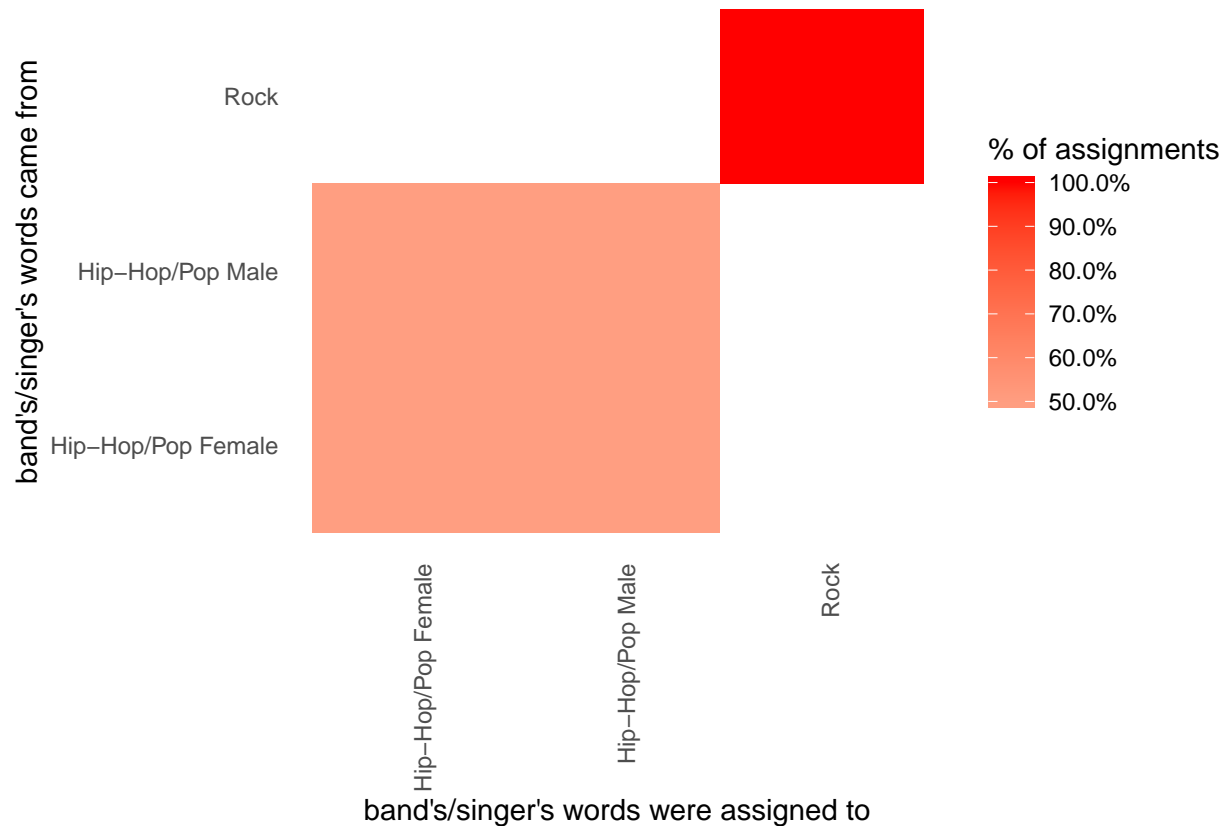
```
## # A tibble: 23,150 x 6
##   document.x      term count .topic document.y      gamma
##   <chr>          <chr> <dbl> <dbl> <chr>      <dbl>
## 1 Hip-Hop/Pop Female 0         4      3 Hip-Hop/Pop Female 1.00
## 2 Hip-Hop/Pop Female 0         4      3 Hip-Hop/Pop Male 1.00
## 3 Hip-Hop/Pop Female 1        66      3 Hip-Hop/Pop Female 1.00
## 4 Hip-Hop/Pop Female 1        66      3 Hip-Hop/Pop Male 1.00
## 5 Hip-Hop/Pop Male 1         29      3 Hip-Hop/Pop Female 1.00
## 6 Hip-Hop/Pop Male 1         29      3 Hip-Hop/Pop Male 1.00
## 7 Hip-Hop/Pop Female 1'2         2      3 Hip-Hop/Pop Female 1.00
## 8 Hip-Hop/Pop Female 1'2         2      3 Hip-Hop/Pop Male 1.00
## 9 Hip-Hop/Pop Female 10         2      3 Hip-Hop/Pop Female 1.00
## 10 Hip-Hop/Pop Female 10         2      3 Hip-Hop/Pop Male 1.00
## # ... with 23,140 more rows
```

*#here document.x is the "true" band whereas document.y is the "consensus" band.*

#confusion matrix for LDA where each genre is a document

```
library(scales)

genre_assignments %>%
  count(document.x, document.y, wt = count) %>%
  group_by(document.x) %>%
  mutate(percent = n / sum(n)) %>%
  ggplot(aes(document.y, document.x, fill = percent)) +
  geom_tile() +
  scale_fill_gradient2(high = "red", label = percent_format()) +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1),
        panel.grid = element_blank()) +
  labs(x = "band's/singer's words were assigned to",
       y = "band's/singer's words came from",
       fill = "% of assignments")
```



*#male and female artist's lyrics are mistaken as each other's.  
#Rock lyrics aren't mistaken for another genre.*

#Wordclouds

```
library(wordcloud)
```

```
## Warning: package 'wordcloud' was built under R version 3.6.3
```

```
## Loading required package: RColorBrewer
```

```
library(wordcloud2)
```

```
## Warning: package 'wordcloud2' was built under R version 3.6.3
```

```
library(RColorBrewer)
```

```
wordcloud_df <- words %>% unique() %>% select(word, n)
```

```
set.seed(1234)
```

```
wordcloud(words = wordcloud_df$word, freq = wordcloud_df$n, min.freq = 1,  
          max.words = 100, random.order = FALSE, rot.per = 0.35, colors = brewer.pal(8, "Dark2"))
```

